

Part 1

Question 1

```
df <- read.csv('/Users/jazzopardi/Desktop/R/699/tech_salary_data.csv')
```

```
df['level'] = NULL # file description says we won't be using this for analysis, so removing from df
```

Question 2

```
str(df)
```

```
'data.frame': 62642 obs. of 23 variables:
 $ timestamp      : chr "6/7/2017 11:33:27" "6/10/2017 17:11:29" "6/11/2017 14:53:57"
 "6/17/2017 0:23:14" ...
 $ company        : chr "Oracle" "eBay" "Amazon" "Apple" ...
 $ title          : chr "Product Manager" "Software Engineer" "Product Manager" "Software
 Engineering Manager" ...
 $ totalyearlycompensation: int 127000 100000 310000 372000 157000 208000 300000
 156000 120000 201000 ...
 $ location       : chr "Redwood City, CA" "San Francisco, CA" "Seattle, WA" "Sunnyvale,
 CA" ...
 $ yearsofexperience : num 1.5 5 8 7 5 8.5 15 4 3 12 ...
 $ yearsatcompany   : num 1.5 3 0 5 3 8.5 11 4 1 6 ...
 $ tag            : chr NA NA NA NA ...
 $ basesalary      : int 107000 0 155000 157000 0 0 180000 135000 0 157000 ...
 $ stockgrantvalue  : num 20000 0 0 180000 0 0 65000 8000 0 26000 ...
 $ bonus           : num 10000 0 0 35000 0 0 55000 13000 0 28000 ...
 $ gender          : chr NA NA NA NA ...
 $ otherdetails     : chr NA NA NA NA ...
 $ cityid          : int 7392 7419 11527 7472 7322 11527 11521 11527 11521 11527 ...
 $ dmaid           : int 807 807 819 807 807 819 819 819 819 819 ...
 $ rowNumber       : int 1 2 3 7 9 11 12 13 15 16 ...
 $ Masters_Degree   : int 0 0 0 0 0 0 0 0 0 0 ...
 $ Bachelors_Degree : int 0 0 0 0 0 0 0 0 0 0 ...
 $ Doctorate_Degree : int 0 0 0 0 0 0 0 0 0 0 ...
 $ Highschool       : int 0 0 0 0 0 0 0 0 0 0 ...
 $ Some_College     : int 0 0 0 0 0 0 0 0 0 0 ...
 $ Race            : chr NA NA NA NA ...
 $ Education       : chr NA NA NA NA ...
 >
```

Timestamp - categorical, company - categorical, level - categorical, title - categorical,
totalcompensation - numerical, location - numerical, years of exp. + years at company -

numerical, tag - categorical, baselary + grantvalue + bonus - numerical, gender - categorical, otherdetails, categorical, cityid - categorical, dmaid - categorical, rownumber - numerical master_degree - categorical, doctoral_degree - categorical, highschool - categorical, some_college - categorical, race - categorical, education - categorical

Question 3

```
set.seed(10)
```

```
0.6 * nrow(df) # 37585 is 60% (for training)
```

```
sampler <- sample_n(df, nrow(df))
```

```
train.df <- slice(sampler, 1:37585) # we have selected our data for training
```


```
valid.df <- slice(sampler, 37586: nrow(df)) # we have 40% (for testing)
```

Data partitioning is important because we need to keep a subset of the available data so that we can use it to verify the model we create with the training data.

As such, the data we don't use is appropriately called testing data. In an ideal situation, you will have training, validation and testing data - but that can be expensive and time consuming.

Question 4

```
ggplot(train.df, aes(x= yearsofexperience,  
  y= totalyearlycompensation)) + geom_point(  
  size = 2, shape = 23) + geom_smooth(method = lm)
```



According to the best line fit, as years of experience increase, so does yearly compensation. This makes sense given that the more experience you have, the more knowledge and expertise you accumulate, and thus the more likely you are to be better compensated. It is important to note that there are several outliers in this training data.

Question 5

```
names(train.df)
```

```
cor.data <- train.df[, c('yearsofexperience', 'totalyearlycompensation')]
```

```
cor(cor.data) # correlation matrix = 0.4135291
```

```
cor.test(train.df$yearsofexperience, train.df$totalyearlycompensation)
```

	yearsofexperience	totalyearlycompensation
yearsofexperience	1.0000000	0.4135291
totalyearlycompensation	0.4135291	1.0000000

Pearson's product-moment correlation

data: train.df\$yearsofexperience and train.df\$totalyearlycompensation

t = 88.049, df = 37583, p-value < 0.000000000000000022

alternative hypothesis: true correlation is not equal to 0

95 percent confidence interval:

0.4051130 0.4218752

sample estimates:

cor

0.4135291

The p value is close to 0, which indicates that the relationship is statistically significant. The correlation sits at 0.413, which suggests a low/medium correlation between the two variables.

Question 6

```
model <- lm(totalyearlycompensation ~ yearsofexperience, data = train.df)
```

```
summary(model)
```

Call:

```
lm(formula = totalyearlycompensation ~ yearsofexperience, data = train.df)
```

Residuals:

```
    Min      1Q  Median      3Q     Max
-718441 -71405 -14208  45121 4569490
```

Coefficients:

```
              Estimate Std. Error t value      Pr(>|t|)
(Intercept)    145879      1028  141.83 <0.0000000000000002 ***
yearsofexperience  9776       111   88.05 <0.0000000000000002 ***
```

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 125600 on 37583 degrees of freedom

Multiple R-squared: 0.171, Adjusted R-squared: 0.171

F-statistic: 7753 on 1 and 37583 DF, p-value: < 0.00000000000000022

Question 7

From the summary() function, we can deduce that the minimum residual is -718441 and the maximum residual is 4569490.

```
train.df$resid <- model$residuals
```

```
View(train.df)
```

For the highest residual value, the observation earned \$4,950,000.

For the lowest residual value, the observation earned \$102,000.

```
View(train.df) # highest and lowest residuals
```

```
max(model$fitted.values)
```

```
4950000 - 380509 # actual vs predicted
```

```
102000 - 820441.3 # actual vs predicted
```

The residuals are calculated by subtracting the model's predicted values from actual value.

There are more factors to take into account when calculating compensation, including company size, exogenous factors to the economy, the observation's own career goals etc. - SLR does not

work for an analysis like this. What's worth noting about these results is that for the lowest residual, the yearsatcompany is listed as 69 - which is most likely a human error in the dataset.

Question 8

```
summary(model)
```

Call:

```
lm(formula = totalyearlycompensation ~ yearsofexperience, data = train.df)
```

Residuals:

```
    Min      1Q  Median      3Q     Max
-718441 -71405 -14208  45121 4569490
```

Coefficients:

```
              Estimate Std. Error t value      Pr(>|t|)
(Intercept)    145879      1028   141.83 <0.0000000000000002 ***
yearsofexperience  9776        111   88.05 <0.0000000000000002 ***
```

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 125600 on 37583 degrees of freedom

Multiple R-squared: 0.171, Adjusted R-squared: 0.171

F-statistic: 7753 on 1 and 37583 DF, p-value: < 0.00000000000000022

The regression equation is $y = 145879 + 9776(x)$

Input year = 15

```
res <- 145879 + (9776*15)
```

```
res = 292519
```

Question 9

```
pred <- predict(model, train.df)
```

```
accuracy(pred, train.df$totalyearlycompensation)
```

```
pred_two <- predict(model, valid.df)
```

```
accuracy(pred_two, valid.df$totalyearlycompensation)
```

	ME	RMSE	MAE	MPE	MAPE
Test set	0.000000001918558	125636.1	83081.14	-40.51568	62.15659

	ME	RMSE	MAE	MPE	MAPE
Test set	-21.36998	124265.8	82168.27	-38.63684	59.99024

The purpose of this comparison is to see how well our model works on data it hasn't seen before. This is important for predictive modeling because future data will be unknown.

The MAE in both models are similar but large, suggesting that this model may not be the most accurate. Similarly the standard deviation between the residuals (the RMSE) is also large, again suggesting that this isn't a good fit. The culprit here is the outliers, which distort our model significantly given that they are considerable outliers. Also, it's a SLR model, which is ineffective given that there are undoubtedly more factors to take into account when predicting total yearly compensation.

Question 10

```
sd(train.df$totalyearlycompensation)
```

137989

```
summary(model)
```

Residual standard error: 125600 on 37583 degrees of freedom

There is a difference of 12,389 between the standard deviation of the overall model and the standard deviation of the residuals.

They are similar, all things considered, suggesting that the model fit well according to the data given - despite there being outliers.

Part 2

```
train.df.two <- read.csv('/Users/jazzopardi/Desktop/R/699/tech_salary_data.csv')
```

Question 1

```
anyNA(train.df.two) # there are NA values
```

```
which(colSums(is.na(train.df.two)) > 0)
```

```
# level # tag # gender # otherdetails # dmaid # race # education
```

These columns have missing values:

level	tag	gender	otherdetails	dmaid	Race	Education
3	9	13	14	16	23	24

b) NA values can produce biased estimates and lead to invalid conclusions.

c)

```
perc <- colSums(is.na(train.df.two))
```

```
x <- sort(perc, decreasing = TRUE)
```

```
names(x[(x/nrow(train.df.two)*100) > 20]) # gender, other details, race, education
```

```
train.df.two <- subset(train.df.two, select = -c(gender, otherdetails, Race, Education))
```

```

> perc <- colSums(is.na(train.df.two))
> perc
      timestamp      company      level      title totalyearlycompensation
      0          0          15          0          0
      location  yearsofexperience  yearsatcompany      tag      basesalary
      0          0          0          808          0
      stockgrantvalue      bonus      gender      otherdetails      cityid
      0          0      19540      22504          0
      dmaid      rowNumber  Masters_Degree  Bachelors_Degree  Doctorate_Degree
      2          0          0          0          0
      Highschool      Some_College      Race      Education
      0          0      40215      32272
> x <- sort(perc, decreasing = TRUE)
> x
      Race      Education      otherdetails      gender      tag
      40215      32272      22504      19540      808
      level      dmaid      timestamp      company      title
      15          2          0          0          0
totalyearlycompensation      location      yearsofexperience      yearsatcompany      basesalary
      0          0          0          0          0
      stockgrantvalue      bonus      cityid      rowNumber      Masters_Degree
      0          0          0          0          0
      Bachelors_Degree      Doctorate_Degree      Highschool      Some_College
      0          0          0          0
> names(x[(x/nrow(train.df.two)*100) > 20])# gender, other details, race, education
[1] "Race"      "Education" "otherdetails" "gender"
> train.df.two <- subset(train.df.two, select = -c(gender, otherdetails, Race, Education))
>

```

d)

```
names(x[(x/nrow(train.df.two)*100) < 20 & (x/nrow(train.df.two)*100) > 0]) # tag # dmaid
```

```
train.df.two$tag[is.na(train.df.two$tag)] <- sort(table(train.df.two$tag), decreasing = TRUE)[1]
```

```
train.df.two$dmaid[is.na(train.df.two$dmaid)] <- sort(table(train.df.two$dmaid), decreasing = TRUE)[1]
```

e)

```
str(train.df.two)
```

```
# rowid is categorical
```

```
train.df.two <- subset(train.df.two, select = -c(rowNumber)) # removing row number
```

A categorical variable with unique values and levels won't be good for predicting because the unique values may rarely occur in real life and thus their impact on the model would be negligent. In this case, it's better to remove the variable altogether or group it according to predefined levels.

f) The three location-related variables are:

1) Location `1050 Levels: Aachen, NW, Germany Aarhus, AR, Denmark Aberdeen Proving Ground, MD Abingdon, MD`

2) CityID `1045 Levels: 0 10 1153 1180 1182 1205 1206 1211 1221 1222`

3) Dmaid `50 Levels: 0 500 501 502 503 504`

```
train.df.two <- train.df.two[train.df.two$dmaid %in% names(sort(table(train.df.two$dmaid), decreasing = TRUE)[1:8]),]
```

g)

```
check <- as.data.frame(sort(table(train.df.two$company), decreasing=TRUE)[1:8])
```

```
check <- c('Amazon','Microsoft','Google','Facebook','Apple','Oracle',  
          'Salesforce', 'Cisco')
```

```
saldat3 <- filter(train.df.two, company %in% check)
```

```
nrow(saldat3)
```

25781

h)

```
check_two <- as.data.frame(sort(table(train.df.two$title), decreasing=TRUE)[1:8])
```



```
check_two <- c('Software Engineer', 'Product Manager','Software Engineering Manager',  
              'Data Scientist','Hardware Engineer','Product Designer','Technical Program Manager',  
              'Solution Architect')
```

```
saldat3 <- filter(saldat3, title %in% check_two)
```

```
nrow(saldat3) # 24391
```

i)

```
check_three <- as.data.frame(sort(table(train.df.two$tag), decreasing=TRUE)[1:8])
```

```
check_three<- c('Distributed Systems (Back-End)', 'Full Stack','API Development (Back-End)',  
               'ML / AI','Web Development (Front-End)','Product','Data',  
               'DevOps')
```

```
saldat3 <- filter(saldat3, tag %in% check_three)
```

```
nrow(saldat3) # 16411
```

j)

```
saldat3$level <- NULL  
saldat3$timestamp <- NULL
```

k)

```
saldat3$basesalary <- NULL  
saldat3$bonus <- NULL  
saldat3$stockgrantvalue <- NULL
```

Question 2

```
set.seed(10)
```

```
0.6 * nrow(saldat3) # 9450
```

```
sampler <- sample_n(saldat3, nrow(saldat3))
```

```
new.train.df <- slice(sampler, 1:9450) # we have selected our data for training
```

```
new.valid.df <- slice(sampler, 9451: nrow(saldat3)) # we have 40% for testing
```

Question 3

```
cor.data <- new.train.df[, c('totalyearlycompensation', 'yearsofexperience',  
                             'yearsatcompany')]
```

```
cor(cor.data)
```

```
               totalyearlycompensation yearsofexperience yearsatcompany  
totalyearlycompensation      1.0000000      0.4865041      0.2734783  
yearsofexperience            0.4865041      1.0000000      0.5167999  
yearsatcompany              0.2734783      0.5167999      1.0000000
```

Question 4

```
salary.lm <- lm(totalyearlycompensation ~ ., data = new.train.df)
```

```
salary.lm.step <- step(salary.lm, direction = 'backward')
```

```
summary(salary.lm.step)
```

```
> salary.lm.step <- step(salary.lm, direction = 'backward')  
Start: AIC=219055.1  
totalyearlycompensation ~ company + title + yearsofexperience +  
  yearsatcompany + tag + dmaid + rowNumber + Masters_Degree +  
  Bachelors_Degree + Doctorate_Degree + Highschool + Some_College
```

```
Step: AIC=219053.2  
totalyearlycompensation ~ company + title + yearsofexperience +  
  yearsatcompany + tag + dmaid + rowNumber + Masters_Degree +  
  Bachelors_Degree + Doctorate_Degree
```

	Df	Sum of Sq	RSS	AIC
<none>			109461687680538	219053
- Masters_Degree	1	30467246568	109492154927106	219054
- Bachelors_Degree	1	116352489345	109578040169883	219061
- yearsatcompany	1	222150066131	109683837746669	219070
- rowNumber	1	482227536474	109943915217012	219093
- Doctorate_Degree	1	533280298015	109994967978553	219097
- tag	7	2003723900866	111465411581404	219211
- title	7	5103524222195	114565211902732	219470
- company	7	16069742910624	125531430591162	220334
- dmaid	7	16827352714643	126289040395181	220391
- yearsofexperience	1	23140568952657	132602256633194	220864

```
> summary(salary.lm.step)  
  
Call:  
lm(formula = totalyearlycompensation ~ company + title + yearsofexperience +  
  yearsatcompany + tag + dmaid + rowNumber + Masters_Degree +  
  Bachelors_Degree + Doctorate_Degree, data = new.train.df)  
  
Residuals:  
    Min       1Q   Median       3Q      Max  
-420073  -50095  -7802    33608  4466329  
  
Coefficients:  
              Estimate Std. Error t value Pr(>|t|)  
(Intercept)    5563.8908   8589.4121   0.648    0.51715  
companyApple    31161.9833   5849.4094   5.327 0.000000101944625 ***  
companyCisco   -57561.9101   7787.4495  -7.392 0.000000000000157 ***  
companyFacebook 116492.3762   4093.7275  28.456 < 0.0000000000000002 ***  
companyGoogle   55262.8401   3803.2656  14.530 < 0.0000000000000002 ***  
companyMicrosoft -28566.5178   3136.9814  -9.106 < 0.0000000000000002 ***  
companyOracle  -29171.7317   5455.3337  -5.347 0.0000000091328639 ***  
companySalesforce 10650.9112   5858.1447   1.818    0.06907 .  
titleHardware Engineer -29551.3076  27124.5453  -1.089    0.27598  
titleProduct Designer -15993.3318  18094.0307  -0.884    0.37677
```

```

tagFull Stack          217.6260    3729.7339    0.058          0.95347
tagML / AI             37854.9122    4596.5154    8.236 < 0.0000000000000002 ***
tagProduct             24135.0490    8964.1917    2.692          0.00711 **
tagWeb Development (Front-End) 541.1245    5732.0459    0.094          0.92479
dmaid501              120765.0051    5773.0826   20.919 < 0.0000000000000002 ***
dmaid506              100225.9500    8595.7786   11.660 < 0.0000000000000002 ***
dmaid511               80625.2104    9763.7606    8.258 < 0.0000000000000002 ***
dmaid635               72008.7805   10299.8213    6.991    0.000000000002909 ***
dmaid803               87661.5883   10061.7447    8.712 < 0.0000000000000002 ***
dmaid807              134830.1864    3949.8008   34.136 < 0.0000000000000002 ***
dmaid819              119648.3539    3552.3490   33.681 < 0.0000000000000002 ***
rowNumber              0.3600        0.0559     6.440    0.000000000125078 ***
Masters_Degree         -4712.0576    2910.8169   -1.619          0.10552
Bachelors_Degree       -11471.9996    3626.3706   -3.163          0.00156 **
Doctorate_Degree       43798.5849    6467.0053    6.773    0.000000000013403 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 107800 on 9415 degrees of freedom
Multiple R-squared:  0.4932,    Adjusted R-squared:  0.4914
F-statistic: 269.5 on 34 and 9415 DF,  p-value: < 0.0000000000000022
> |

```

Question 5

```

sst <- new.train.df$totalyearlycompensation - mean(new.train.df$totalyearlycompensation)
sst <- sst^2
sst <- sum(sst) # the difference between the mean and the observed value

```

215998347738862

Question 6

```

ssr <- salary.lm.step$fitted.values - mean(new.train.df$totalyearlycompensation)

ssr <- ssr^2

```

```
ssr <- sum(ssr)
```

106536660058325


Question 7

```
ssr / sst
```

0.4932291

Question 8

The predictor I chose with “Master’s Degree.”



89.5 % of the curve is shaded

$$1 - 0.895 = 0.105 \text{ (p value)}$$

The p-value is the “unshaded” part of the curve which tells us whether the value chosen is up to chance or not.

Masters_Degree	-4712.0576	2910.8169	-1.619	0.10552
----------------	------------	-----------	--------	---------

Question 10

```
newframe <- data.frame(company = 'Facebook',  
  title = 'Product Designer',  
  yearsofexperience = 5,  
  yearsatcompany = 2,  
  tag = 'Data',  
  dmaid = '501',  
  Masters_Degree = 1,  
  Bachelors_Degree = 1,  
  Doctorate_Degree = 0,  
  rowNumber = 0)
```

```
predict(salary.lm.step, newframe)
```

The tech worker will make \$269888 annually based on this model.

```
> predict(salary.lm.step, newframe)  
1  
269888.6
```

Question 11

```
pred <- predict(salary.lm.step, new.train.df)  
accuracy(pred, new.train.df$totalyearlycompensation)
```

```
pred_two <- predict(salary.lm.step, new.valid.df)  
accuracy(pred_two, new.valid.df$totalyearlycompensation)
```

```

> accuracy(pred, new.train.df$totalyearlycompensation)
               ME      RMSE      MAE      MPE      MAPE
test set 0.000000009145604 107867.1 62149.71 -9.531122 27.81096
> pred_two <- predict(salary.lm.step, new.valid.df)
> accuracy(pred_two, new.valid.df$totalyearlycompensation)
               ME      RMSE      MAE      MPE      MAPE
test set 1554.252 121618 60921.05 -8.394922 26.6268

```

We can see here that the MAE for the MLR is lower than the MAE for SLR which suggests that our MLR model is more accurate. This makes sense given that the conditions of SLR rarely exist in everyday life.

The MAE in the train and testing set are similar which means our model is good at predicting based on unknown data. Of interesting note here is that the ME for the testing set is very small - perhaps a sign of overfitting.